

Grow My Money

An Architecture for Deterministically Authorized AI Financial Agents

From recommendation to bounded agency: permission-aware context, explainable opportunity discovery, deterministic risk, policy-gated execution, canonical ledger truth, and auditable evidence.

THE CONTROL DOCTRINE

AI may discover and propose. Policy determines permission. Risk determines exposure. The user approves where required. Controlled financial infrastructure executes. The ledger establishes truth.

Technical / Economic White Paper

September 2026 | SunRey

Conceptual architecture. Regulated capabilities require jurisdiction-specific licensing, service-provider coverage, legal review, model validation, security approval, and production activation. No investment return is guaranteed.

Executive Abstract

Grow My Money is the operating brain behind SunRey's Personal Economic Operating System. Its design starts from a broader problem than conventional robo-advice: not merely how to allocate an investment portfolio, but what legitimate, authorized action could improve a person's total economic position next.

The architecture therefore treats AI as an intelligence and proposal layer, not as an unrestricted financial principal. The system can inspect permissioned economic context, discover opportunities across cash, foreign exchange, investing, rewards, subscriptions, commerce, research and other domains, and assemble a ranked Growth Plan. Consequential actions become ActionIntents. They remain proposals until deterministic control systems establish identity, authority, jurisdiction, compliance, risk, purpose, model state and consent, with human approval where required.

Within the investment domain, SunRey separates specialized intelligence into an Agentic Capital Mesh: macro/FX, relative-value, volatility, microstructure, execution and meta-allocation agents. These agents may compete and cooperate, but a deterministic Risk Engine outranks them on exposure, leverage, liquidity, concentration, drawdown and kill switches.

The result is a distinct architecture for agentic finance: increase the breadth and speed of machine intelligence while reducing ambiguity about who or what is allowed to execute. Financial autonomy becomes bounded by explicit mandates, machine-readable legal capability, deterministic controls, signed execution authority, canonical accounting and reconstructable evidence.

CENTRAL CLAIM

The differentiator is not that AI can generate financial advice. It is that AI-generated economic action can be translated into a deterministic authorization and evidence chain before execution.

Why AI Financial Advice Is the Wrong Frame

The phrase "AI financial advice" describes a conversational interface, but it does not define a safe operating architecture. A useful financial agent may need to reason across cash reserves, upcoming obligations, currency exposure, portfolio risk, rewards, spending leakage, data-market opportunities and user goals. Once the system can propose consequential actions across those domains, the design problem becomes one of authority, not only intelligence.

SunRey addresses this by separating three questions that are often collapsed: What does the system know? What does it recommend? What is permitted to happen? The first belongs to permission-aware economic context; the second to AI and optimization; the third to deterministic policy, risk, user mandate and execution controls.


- A model can be wrong while a control system remains correct.
- A recommendation can be economically attractive yet legally unavailable in a customer's jurisdiction.
- An investment can fit a market thesis while violating liquidity, concentration or drawdown limits.
- A data point can be known while remaining unauthorized for a particular purpose or model.
- An action can pass policy and risk checks yet still require explicit user approval before execution.

DESIGN SHIFT

The objective is not autonomous finance without friction. The objective is high-agency financial intelligence with explicit, inspectable authority boundaries.

System Thesis: Intelligence Around a Control Spine

Grow My Money - Deterministically Authorized Control Loop



Core invariant: AI proposes. Policy permits. Risk constrains. Authorized infrastructure executes. The ledger records.

No AI model, agent, recommendation service, or market-research component is itself financial execution authority.

The SFF 2.0 architecture is built around a control spine inherited from the Sovereign Financial Fabric: ActionIntents, deterministic policy evaluation, risk controls, signed Execution Authority, segregated financial domains, ledgers and evidence. Around that spine, SunRey adds economic intelligence capable of understanding the customer, searching for opportunities and coordinating specialized agents.

This design preserves one key invariant: no AI model gains execution authority merely because it is intelligent, personalized, confident or profitable in simulation. The proposal must survive an independent authorization path.

The Five First-Class Systems

System	Question it owns	Primary output
Personal Economic Graph	What do we know, and what may be used for this decision?	Permission-aware economic context
Growth Orchestrator	What legitimate actions could improve the user's position?	Ranked Growth Plan and ActionIntents
Agentic Capital Mesh	How should authorized investment risk be deployed?	Strategy proposals and execution plans
Personal Economic Value Engine	What is a specific authorized contribution worth in context?	Contribution valuation and compensation basis
Regulatory Digital Twin	What is SunRey legally and operationally capable of doing here and now?	Machine-readable capability and control state

These are treated as first-class infrastructure rather than helper logic. Each is expected to own domain models, APIs, events, storage, security boundaries, permission checks, observability, evidence, tests, documentation and model/version governance. The practical consequence is that the core concepts become enforceable in code rather than aspirational product language.

SFF 2.0

The original financial fabric controls what may happen. The five first-class systems add the intelligence required to understand context, discover opportunity, deploy specialized strategies, value authorized contribution and map legal capability.

Personal Economic Graph

The Personal Economic Graph (PEG) is a permission-aware model of the user's economic world. It may connect accounts, income, expenses, assets, debts, goals, subscriptions, currencies, rewards, interests and other authorized context into relationships that agents can reason over.

Its central question is not simply "what data exists?" but "what information is relevant and permitted for this decision?" That distinction is essential in an agentic system because broad access to data can otherwise be mistaken for broad authority to use it.

Relationship metadata	Why it matters
Source and verification	Separates observed, asserted, derived and externally verified facts.
Sensitivity	Controls handling of financial, behavioral or other sensitive context.
Purpose	Binds information use to a declared economic task.
Consent and expiry	Makes authorization revocable and time-bounded.
Jurisdiction	Allows data-use and product rules to vary by applicable legal context.
Permitted agent / model	Prevents every model from automatically inheriting every data permission.

SAFETY PRINCIPLE

Knowing something never implies permission to use it. Consent and purpose remain authoritative.

From Context to Deployable Insight

The graph is useful because it can turn fragmented facts into decision context without itself becoming financial authority. For example, the source architecture illustrates a user with monthly income of \$15,000, ordinary obligations of \$8,000, a required liquidity reserve of \$40,000 and current cash of \$52,000. The graph can support the inference that roughly \$12,000 may be deployable - but only subject to the user's mandate, policy and risk.

This pattern generalizes. The graph may expose a mismatch between idle cash and liquidity needs, recurring fees and user priorities, currency holdings and expected obligations, or portfolio concentration and an approved risk budget. The graph supplies context; it does not authorize a trade, transfer or account change.

- Facts can be verified without becoming executable instructions.
- Derived insights should retain lineage to the facts and permissions that support them.
- Agent outputs should identify the graph facts materially influencing a recommendation.
- Expired or revoked permissions should remove data from future decision contexts without rewriting historical evidence.

Customer Mandates: The Human Boundary

A financial agent needs more than a profile. It needs an explicit mandate describing what the user wants and what the system may consider. The SunRey design treats mandate generation as a control boundary between user intent and machine action.

Mandate dimension	Illustrative meaning
Goal	Increase cash yield, reduce fees, build reserves, invest for a target, reduce currency risk.
Time horizon	Immediate, monthly, multi-year or goal-specific.
Liquidity floor	Cash that must remain readily available.
Risk boundaries	Permitted volatility, drawdown, leverage, concentration and asset classes.
Action scope	What may be proposed, auto-prepared or require explicit approval.
Data scope	What information may be used for the task and by which models.
Frequency / duration	How often the mandate may act and when it expires or must be renewed.

Mandates should be versioned and reconstructable. A later audit should be able to identify the exact user instruction, risk boundaries and permitted scope that existed when an ActionIntent was created.

Growth Orchestrator

The Growth Orchestrator is the operating brain behind the Grow My Money control. It continuously gathers candidate opportunities, evaluates them against the Personal Economic Graph and user mandates, and assembles an explainable Growth Plan.

Opportunity family	Representative search space
Cash optimization	Yield, reserve placement, liquidity routing.
FX	Permitted routes, spreads, fees, timing and currency needs.
Investing	Portfolio changes, validated strategies, risk-budget deployment.
Rewards / merchant offers	Cashback, rewards and negotiated value.
Subscriptions / bills	Cancellation, plan change and verified cost reduction.
Research / data opportunities	Permissioned participation where legally and operationally allowed.
Commerce optimization	Price, coupon, reward and payment-method combinations.

The Orchestrator does not execute. Its primary output is a ranked Growth Plan decomposed into ActionIntents. This ensures the optimization layer cannot bypass authorization simply by embedding execution inside a recommendation service.

ActionIntents: The Boundary Object

ActionIntents are the architectural boundary between probabilistic intelligence and deterministic execution. Any consequential action proposed by an AI agent is serialized into a structured object that can be evaluated independently of the model that created it.

ActionIntent field family	Purpose
Actor / user / mandate	Establish whose objective is being served and under what instruction.
Action type and parameters	Describe the exact transfer, order, rebalance, cancellation or other action.
Economic rationale	Provide the agent's reason and expected benefit without granting authority.
Inputs and provenance	Identify material data, model and market inputs.
Risk estimate	Expose intended exposure and expected constraints.
Jurisdiction / product context	Allow policy evaluation against the correct capability state.
Approval requirement	State whether explicit user or manual approval is required.
Idempotency / expiry	Prevent stale or duplicate financial effects.

BOUNDARY RULE

An ActionIntent is a proposal. It is not a transaction, an order, a transfer, an authorization or a ledger posting.

Proof of Permission

The Compliance Kernel evaluates whether a consequential ActionIntent can progress. The source architecture defines six foundational proofs and adds Model State and Consent as explicit control dimensions for AI- and data-intensive actions.

Proof / control	Question
Identity	Who is involved and is identity current and verified?
Authority	Who authorized this exact action and within what limits?
Jurisdiction	Which law, regulator, legal entity and regional operating context govern?
Compliance	Is the product or activity permitted given KYC/AML/sanctions/product rules?
Risk	Is financial exposure within deterministic limits?
Purpose	Is sensitive information used only for an authorized purpose?
Model State	Is the economically consequential model approved and permitted for this product?
Consent	Does a current granular consent object cover the requested use?

Deterministic Outcomes

The control layer is more expressive than a binary allow/block switch. The source architecture defines deterministic outcomes that can condition, defer or escalate an action without granting the model arbitrary discretion.

Outcome	Operational meaning
ALLOW	Proceed to the next controlled stage.
ALLOW WITH CONDITIONS	Proceed only with explicit constraints.
STEP-UP AUTHENTICATION	Require stronger user authentication.
REQUIRE DISCLOSURE	Present required information before approval or execution.
REQUIRE KYC REFRESH	Update identity/compliance state.
REQUIRE SOURCE-OF-FUNDS EVIDENCE	Collect additional evidence before continuing.
REQUIRE MANUAL REVIEW	Route to an authorized human control function.
DEFER	Do not execute now; await a future state or decision.
BLOCK	Do not permit the action.

Every outcome, including a block, can be sealed into the Evidence Vault with identity state, policy version, model version, jurisdiction, purpose, risk result and final decision. This allows the system to explain not only what executed, but also what was prevented and why.

Three-Plane Authority Separation

Three-Plane Authority Separation

1. INTELLIGENCE PLANE

Graph reasoning, market research, opportunity discovery, agent proposals, model inference.

2. CONTROL PLANE

Mandates, consent, jurisdiction, compliance, deterministic risk, approvals, capability state.

3. EXECUTION + TRUTH PLANE

Signed execution authority, controlled connectors, canonical ledger, reconciliation, evidence.

Design objective: useful AI without implicit authority.

This separation is the core of the architecture. The intelligence plane may use probabilistic reasoning. The control plane turns user mandates, policy, regulatory capability and risk into deterministic constraints. The execution and truth plane performs only authorized actions and records canonical outcomes.


WHY IT MATTERS

A financial AI can be made more capable without simultaneously making it more powerful in the legal or operational sense. Capability and authority are deliberately decoupled.

Agentic Capital Mesh

Agentic Capital Mesh

Specialized intelligence beneath a deterministic Risk Engine



Within the investment domain, SunRey avoids treating one monolithic model as trader, portfolio manager, risk officer and execution engine at the same time. Specialist agents compete and cooperate, while the Risk Engine remains the superior control authority.

Macro / FX and Relative-Value Agents

Macro/FX agents focus on regimes, rates, currencies and cross-asset conditions. Their useful role is to identify potential exposures or hedges, not to override portfolio limits. Relative-value agents search for short-horizon mispricing or market-neutral structures, but their proposals remain constrained by capacity, liquidity and concentration.

This architecture encourages specialization because the inputs, horizons, validation standards and failure modes differ. A macro signal may be slow-moving and regime-dependent; a relative-value signal may decay rapidly and be sensitive to crowding. Treating these as distinct agent classes allows the system to attach different model governance, data requirements and risk budgets to each.

- Agents produce evidence-rich proposals rather than opaque instructions.
- The same thesis can be rejected if liquidity, concentration or jurisdiction state changes.
- Strategy research should retain replayable inputs and versioned model state.
- Risk and authorization remain independent of the agent's confidence score.

Volatility and Microstructure Agents

Volatility agents reason about volatility surfaces, options structures and convexity. Their control boundary is explicit: Greeks, tail-risk limits and other deterministic constraints remain enforced independently of the agent. Microstructure agents reason about order-book state, spread, queue position and adverse selection, but have no direct venue access without authorization.

This matters because the closer an agent operates to market execution, the more dangerous it becomes to confuse intelligence with authority. A high-frequency or microstructure model can generate recommendations at machine speed, but the architecture still requires a deterministic path for permissions, risk limits, execution policy and kill switches.

NON-NEGOTIABLE

Machine speed does not eliminate policy. It increases the importance of pre-authorized envelopes, deterministic limits and immediate kill capability.

Execution Agent and Meta Allocator

The Execution agent focuses on venue selection, order slicing, timing, slippage and fill quality. It does not own the investment thesis and it does not own the risk budget. Its job is to turn an already permitted exposure into an execution plan consistent with policy and market conditions.

The Meta Allocator allocates risk budget across validated strategies. In the source design, the Risk Engine outranks all allocation decisions. This prevents a portfolio-level optimizer from silently creating leverage or concentration simply because multiple specialized agents independently report attractive opportunities.

Layer	May optimize	May not override
Strategy agent	Signal quality and proposed exposure	Mandate, product eligibility, risk limits
Meta Allocator	Risk-budget distribution across validated strategies	Global risk engine, liquidity floor, concentration
Execution agent	Venue, timing, slicing, expected implementation cost	Authorization, exposure limits, kill switches

Deterministic Risk Engine

The deterministic Risk Engine is the final authority on the dimensions that can threaten the user's capital or the platform's operating integrity. The architecture specifically places exposure, leverage, drawdown, liquidity, concentration and kill switches under deterministic control rather than model discretion.

Risk dimension	Control objective
Exposure	Constrain absolute and relative market exposure.
Leverage	Prevent strategy or portfolio leverage from exceeding approved limits.
Liquidity	Ensure positions can be entered, exited or funded within acceptable constraints.
Concentration	Limit dependence on one asset, strategy, venue or correlated risk source.
Drawdown	Reduce or stop activity when loss thresholds are breached.
Kill switches	Provide immediate deterministic suspension independent of agent state.

Risk is evaluated before execution and should also remain active during position life, execution and settlement. The system should not escalate leverage to chase an optimization target, and the investment component is explicitly allowed to lose money.

Model Lifecycle and Promotion

Specialization only helps if strategies are governed. The SFF 2.0 source explicitly describes a progression through research, replay, shadow, paper and canary stages before controlled production.

Stage	Purpose
Research	Formulate hypotheses, datasets, objectives and baseline comparisons.
Replay	Evaluate behavior against historical or recorded market conditions.
Shadow	Run alongside live or simulated conditions without financial authority.
Paper	Exercise full decision workflows using simulated capital and execution.
Canary	Introduce narrowly bounded production exposure after explicit approval.
Controlled production	Operate only inside approved mandate, risk, model and regulatory capability envelopes.

Model approval is therefore a state that can be checked by the Compliance Kernel. A model can be technically available while remaining disallowed for a product, jurisdiction or customer segment.

Growth Is Broader Than Investment Return

Growth Attribution - Economic Improvement Is Not One Return

SETTLED CASH Deposit yield, realized rewards, settled compensation

UNREALIZED Mark-to-market investment gains/losses

COST AVOIDED FX, subscription, bill, commerce savings

PENDING Approved but unsettled compensation or rewards

The Weekly Economic Delta may aggregate categories for visibility, but must never mislabel the total as investment performance.

Grow My Money is not defined as a single investment strategy. The source architecture activates a customer-scoped portfolio of economic mechanisms, each with its own permission state, risk classification, legal treatment and attribution. This is what allows SunRey to optimize across a person's broader economic life without misrepresenting every benefit as investment performance.

Growth Mechanisms

Mechanism	Value source	Realization treatment
Cash optimization	Deposit yield / cash routing	Settled cash
Investment strategies	Realized and unrealized P&L	Settled cash or unrealized; losses allowed
Research / data opportunities	Permissioned participation	Pending to settled compensation
Merchant offers	Cashback, discounts, rewards	Settled reward or cost avoided
Subscription agent	Cancellation or plan savings	Cost avoided
Bill agent	Provider / plan optimization	Cost avoided
Commerce optimizer	Price, coupon, reward, payment method	Cost avoided or reward
FX optimizer	Reduced spread / fee vs permitted baseline	Cost avoided

Each mechanism can be evaluated independently, then presented together as an economic delta. This preserves a user-centered view of total economic improvement without collapsing accounting categories.

Growth Attribution Ledger

The Growth Attribution Ledger is a separate analytical/accounting layer that explains where economic improvement came from. It does not replace the canonical financial ledger and it does not convert heterogeneous outcomes into a synthetic investment return.

Realization class	Meaning
SETTLED_CASH	Cash or asset value that has actually settled and is available subject to product rules.
UNREALIZED	Mark-to-market gain or loss not yet realized.
COST_AVOIDED	Verified reduction in an expense; not income and not investment return.
PENDING	Expected or approved compensation that has not yet settled.

This distinction becomes a consumer-trust feature. A weekly summary may show economic improvement across several sources, but a negative investment result remains negative; it cannot be hidden inside unrelated savings or rewards.

MEASUREMENT RULE

Never convert total economic improvement into a misleading percentage investment return.

Personal Economic Value Engine

PEVE is part of the five-system architecture because Grow My Money may eventually evaluate permissioned human-information or research opportunities alongside conventional finance. PEVE does not calculate the intrinsic worth of a person. It estimates the market value of a specific, authorized contribution in a specific context.


Relevant signals may include demand, scarcity, freshness, verification quality, completeness, utility, geography, permission scope, privacy sensitivity and compliance cost. The resulting value can inform compensation, but it remains separate from the person's identity, from a social score, and from financial execution authority.

BOUNDARY

PEVE can inform whether an opportunity is economically attractive. Consent, purpose, legal capability and execution controls still determine whether participation is allowed.

Regulatory Digital Twin

Regulatory Digital Twin - Capability as Machine-Readable State



Operational effect: changes update authoritative capability state instead of being scattered through product logic.

Audit effect: the system can reconstruct why an action was permitted, conditioned, deferred, reviewed, or blocked.

The Regulatory Digital Twin (RDT) is a machine-readable model of what SunRey is legally and operationally capable of doing by product, customer, jurisdiction, legal entity, license state, partner, control requirement, policy version and review state. It converts legal capability from scattered assumptions into explicit system state.

The source architecture describes a chain from regulation to requirement to control to policy to software implementation to test to evidence to regulatory report. A product action can therefore be permitted, conditioned or blocked based on current capability state rather than hard-coded country logic inside an AI agent.

Jurisdiction Is a Runtime Constraint

Financial intelligence becomes genuinely operational only when it knows what is available to a particular user in a particular place under a particular operating model. The same proposed action may be permissible in one jurisdiction, require disclosures or partner routing in another, and remain disabled elsewhere.

The RDT enables the Growth Orchestrator to search only within valid capability space. It also gives the system a way to respond to rule changes: update authoritative capability state and its evidence rather than rewriting dozens of agent prompts or product rules.

- A technically implemented capability can remain disabled.
- A provider integration does not itself establish legal permission.
- A rule change can alter the opportunity set without modifying model weights.
- Historical actions can be reconstructed against the policy and legal state that existed at the time.

Execution Authority and Canonical Ledger

After an ActionIntent passes the required controls and approvals, the system issues a signed, short-lived Execution Authority. This artifact - not the AI recommendation - is what authorizes ledger-impacting execution.

The canonical Ledger remains the source of truth for financial movements. The underlying accounting design preserves balanced journals, append-only history, asset segregation, named bridges between asset classes, idempotency and reconciliation. This is crucial because an agentic interface must not create a shadow financial state that diverges from actual accounting.

Ledger invariant	Guarantee
Balance	Every journal balances per asset.
Immutability	Corrections use compensating entries rather than rewriting history.
Authority	Customer-impacting journals require valid execution authority.
Segregation	Deposits, investments, digital assets and corporate assets remain distinct.
Idempotency	Duplicate requests do not duplicate financial effects.
Reconciliation	Internal truth is compared with approved external financial infrastructure.

Evidence Vault and Proof-Carrying Financial Events

A consequential action should be reconstructable as more than a transaction row. The evidence layer can preserve the identity state, user mandate, ActionIntent, model version, policy version, risk result, approval state, execution authority, ledger outcome and relevant external confirmations.

This creates a proof-carrying financial event: the result is linked to the evidence explaining why it was allowed and how it completed. Blocks and deferrals can be evidenced as well. The objective is not only post-hoc logging; it is to make traceability part of the execution architecture.

AUDIT QUESTION

Can SunRey reconstruct exactly what the agent proposed, what data it was permitted to use, which policy and risk limits applied, who approved the action, what executed, and what accounting result followed?

Failure Modes and Safe Degradation

Agentic finance must define what happens when the intelligence layer, external data, policy services or execution infrastructure are unavailable or disagree. The architecture favors fail-closed behavior for authority and graceful degradation for intelligence.

Failure	Expected posture
Model unavailable	Do not invent a recommendation; preserve existing state and allow non-agent functions where safe.
Market data stale	Mark data stale; suppress strategies requiring current state.
Policy service unavailable	Do not execute a consequential action requiring policy approval.
Risk engine unavailable	Block new risk-increasing actions.
Provider execution uncertain	Reconcile before retry; use idempotency to avoid duplicate economic effects.
Consent revoked	Remove affected information from future agent decision contexts.
Jurisdiction capability unresolved	Keep capability disabled or require review.

Security and Authority Boundaries

The highest-value security property is that no single model, provider, operator, database or interface can silently become financial authority. SunRey's broader architecture treats AI as proposal and explanation infrastructure; critical authority remains distributed across identity, permission, compliance, risk, user approval, execution authority and ledger controls.

- Least-privilege capability tokens for agents and models.
- Explicit separation between data access and action authority.
- Model/version registry for economically consequential inference.
- No direct model-generated SQL or uncontrolled access to production financial databases.
- Short-lived signed execution authority rather than persistent broad credentials.
- Kill switches and operating posture flags independent of AI systems.
- Evidence, idempotency and reconciliation around every meaningful economic action.

What the User Experiences

The customer experience can remain simple even though the backend is deliberately compartmentalized. A user may express a goal, review a Growth Plan, inspect explanations and approve bounded actions. The complexity of policy, risk, provider routing, accounting and evidence remains behind the interface without being hidden from audit.

User step	System behavior
Define goal	Compile objective, horizon, liquidity and risk boundaries into a mandate.
Connect permitted context	Build the Personal Economic Graph under source, consent and purpose rules.
Ask Grow My Money	Search eligible opportunity families and rank candidates.
Review Growth Plan	Show expected benefit, risks, assumptions and required approvals.
Approve action	Create/advance the ActionIntent through deterministic controls.
Execute	Use controlled infrastructure only after valid Execution Authority.
Measure	Attribute outcome by source and realization class.
Learn	Use outcomes and approved preferences to improve future proposals without expanding authority automatically.

Illustrative End-to-End Growth Journey

Consider a user whose goal is to improve monthly financial efficiency while maintaining a defined cash reserve and moderate investment risk. The Graph identifies excess cash above the liquidity floor, a recurring subscription no longer aligned with the user's priorities, a future foreign-currency obligation and a portfolio concentration that exceeds the user's target range.

The Growth Orchestrator may produce four candidate actions: move a bounded amount of idle cash to a higher-yield permitted account, cancel or downgrade the subscription, obtain a permitted FX route for the future obligation, and propose a portfolio rebalance. Each candidate becomes an ActionIntent with its own rationale, data lineage, risk profile, product and jurisdiction context.

The subscription change may need only user confirmation; the FX transfer may require corridor-specific policy checks; the portfolio rebalance must pass investment eligibility and deterministic risk; the cash movement must preserve the liquidity floor. Only the actions that receive the required approvals and Execution Authority can affect money. The results are then recorded separately: cash yield as settled cash, subscription savings as cost avoided, FX savings against a permitted baseline, and investment performance as realized or unrealized P&L.

WHY THIS IS DIFFERENT

One consumer command can coordinate multiple economic mechanisms without pretending they are one product, one risk class or one investment return.

Claims SunRey Should Not Make

The source architecture is explicit about several boundaries that should remain visible in any public presentation of Grow My Money.

- Do not guarantee investment returns or promise a fixed weekly growth rate.
- Do not label expense savings, rewards or data compensation as investment performance.
- Do not imply that AI can bypass legal, compliance, risk or user-approval requirements.
- Do not treat software implementation as proof that a regulated capability is legally active.
- Do not let a growth target justify leverage escalation or weakened risk limits.
- Do not collapse deposits, investments, digital assets or corporate assets into one undifferentiated balance or legal category.

These constraints are not merely compliance disclaimers. They are architectural requirements that help define what the system is and is not allowed to do.

Research and Validation Agenda

The value of the architecture should be tested empirically. The most important research questions are not only whether an AI can predict markets, but whether a bounded agentic system can produce measurable economic improvement while maintaining clear authority, reliable accounting and explainable control decisions.

Research question	Possible evaluation
Opportunity quality	Incremental economic value versus non-agent baselines by source.
Authorization safety	Rate of prevented policy/risk violations and false-positive blocks.
Explainability	Whether users and reviewers can trace material factors behind proposals.
Risk discipline	Drawdown, liquidity, concentration and leverage behavior under stress.
Model governance	Performance drift and failure across research/shadow/paper/canary stages.
Attribution integrity	Accuracy of settled, unrealized, cost-avoided and pending classifications.
Regulatory adaptability	Time and error rate required to update capability state after rule changes.

Strategic Differentiation

Most AI-finance products can be described as advice interfaces, research copilots, portfolio tools or automated execution systems. SunRey's architectural thesis is broader: a financial agent should operate across a permissioned model of the user's economic world, coordinate specialized optimization systems, and convert every consequential proposal into a deterministic authorization path.

That makes the core unit of innovation not the chatbot and not the trading model, but the controlled transition from machine inference to economic action. The Personal Economic Graph provides relevant and permitted context. The Growth Orchestrator converts context into a ranked plan. The Agentic Capital Mesh supplies specialist financial intelligence. The Regulatory Digital Twin constrains actions to legal capability. The Compliance Kernel and Risk Engine determine permission and exposure. Execution Authority, the Ledger and Evidence Vault establish what actually happened.

POSITIONING

Grow My Money is best understood as an agentic financial operating architecture - not simply an AI adviser.

Conclusion

Grow My Money proposes a way to make financial agents more capable without making their authority ambiguous. The system can search broadly across the user's authorized economic life, combine deterministic context with probabilistic intelligence, and coordinate multiple specialized agents. But the right to execute remains outside the AI layer.

The architectural pattern is intentionally conservative where money, risk, data rights and legal capability are concerned: explicit mandates, permission-aware context, ActionIntents, deterministic policy, deterministic risk, human approval where required, signed execution authority, canonical ledger truth, attribution and evidence. It is intentionally ambitious in the intelligence layer: continuous opportunity discovery, specialized market agents, multi-domain optimization and machine-readable regulatory capability.

If the approach works, the result is more than automated advice. It is a Personal Economic Operating System in which AI can reason broadly and act usefully while remaining bounded by controls that are inspectable, auditable and independent of the model itself.

AI may discover and propose. Policy determines permission. Risk determines exposure. Controlled infrastructure executes.

Source Basis and Terminology

This white paper synthesizes the current SunRey architecture from the following internal source documents. Historical source documents use earlier Solstice / Pyramid naming; this paper applies the current SunRey product vocabulary while preserving the underlying architecture.

ID	Source	Primary contribution
S1	SFF 2.0 - Five First-Class Architecture Components (August 2026)	Personal Economic Graph, Growth Orchestrator, Agentic Capital Mesh, PEVE and Regulatory Digital Twin.
S2	Comprehensive Product, Architecture, API, Compliance & Growth Systems Report - SFF 2.0 (August 2026)	ActionIntent / Compliance Kernel / Proof of Permission, Grow My Money, risk, ledger, attribution, evidence and regulatory architecture.
S3	SunRey Comprehensive Business Plan (September 2026)	Current SunRey naming, flagship Grow My Money positioning and non-negotiable AI / policy / risk / execution control rule.
S4	Architecture & Feature Reference (August 2026)	Sovereign Financial Fabric, agent runtime, capability tokens, ledger, Growth Attribution Ledger and Evidence Vault.
S5	Master Build Instructions (August 2026)	Non-negotiable financial rules, implementation priorities, simulation-first controls and final product principle.

Selected Terminology

Term	Meaning
ActionIntent	Structured proposal for a consequential economic action; never execution authority by itself.
Execution Authority	Signed, short-lived artifact authorizing a controlled execution after required checks.
Personal Economic Graph	Permission-aware representation of user economic context.
Growth Orchestrator	System that discovers and ranks legitimate opportunities into a Growth Plan.
Agentic Capital Mesh	Portfolio of specialized investment and quantitative agents beneath deterministic risk.
Regulatory Digital Twin	Machine-readable representation of product/jurisdiction/legal capability state.
Growth Attribution Ledger	Layer separating economic improvement by source and realization class.
Evidence Vault	Durable evidence layer supporting reconstruction of decisions and outcomes.

Publication Note

This document describes an architectural thesis and intended system behavior. It is not a representation that every described capability is currently licensed, deployed or available in every jurisdiction. Regulated banking, payments, investment, digital-asset, data-use and other capabilities require appropriate legal entities, service-provider relationships, jurisdiction-specific approvals and operational controls before production activation.

Nothing in this paper is a guarantee of financial performance. Investment outcomes may be negative. Economic improvement from rewards, cost savings, data participation, foreign-exchange optimization or other categories should remain separately classified from investment performance.

The architecture also does not imply that AI models possess legal or financial authority. The defining SunRey design principle is the opposite: AI may research, reason, discover and propose, while permission, risk, approval, execution and accounting truth remain controlled by separate deterministic systems.

SUNREY

Grow My Money - deterministically authorized AI financial agents for a Personal Economic Operating System.